

Homework 1 - Berkeley STAT 157

Handout 1/22/2017, due 1/29/2017 by 4pm in Git by committing to your repository. Please ensure that you add the TA Git account to your repository.

1. Write all code in the notebook.
2. Write all text in the notebook. You can use MathJax to insert math or generic Markdown to insert figures (it's unlikely you'll need the latter).
3. **Execute** the notebook and **save** the results.
4. To be safe, print the notebook as PDF and add it to the repository, too. Your repository should contain two files:
homework1.ipynb and homework1.pdf.

The TA will return the corrected and annotated homework back to you via Git (please give `rythei` access to your repository).

In [3]:

```
from mxnet import ndarray as nd
```

1. Speedtest for vectorization

Your goal is to measure the speed of linear algebra operations for different levels of vectorization. You need to use `wait_to_read()` on the output to ensure that the result is computed completely, since NDArray uses asynchronous computation. Please see http://beta.mxnet.io/api/ndarray/_autogen/mxnet.ndarray.NDArray.wait_to_read.html for details.

1. Construct two matrices A and B with Gaussian random entries of size 4096×4096 .
2. Compute $C = A B$ using matrix-matrix operations and report the time.
3. Compute $C = A B$, treating A as a matrix but computing the result for each column of B one at a time. Report the time.
4. Compute $C = A B$, treating A and B as collections of vectors. Report the time.
5. Bonus question - what changes if you execute this on a GPU?

In [36]:

```
# part 1
A = nd.normal(0, 1, shape = (4096, 4096))
B = nd.normal(0, 1, shape = (4096, 4096))

import time

# part 2
tic = time.time()
C = nd.dot(A,B)
print(C)
print(time.time() - tic)

[[-50.474396  -69.83126  -41.304188  ... -19.597466   35.388355
   78.81215   ]
 [-29.70652   -26.495625   96.28163   ...  63.203274   31.081131
   54.01789   ]
 [ 45.610382   28.23838   -1.6904984 ...  45.466324   -6.7074146
  -47.965813   ]
 ...
 [-81.30978   -46.991093  -49.305668  ... -19.005987   -9.508311
  -24.298445   ]
 [-70.84065   113.742485  -85.99479   ... -54.69081   10.467543
   -7.410914   ]
 [ 26.251398  -87.17         63.974815  ...   2.6107101  25.058563
   50.11858    ]
<NDArray 4096x4096 @cpu(0)>
3.211174964904785
```

In [37]:

```
# part 3
import numpy as np
```

```

import numpy as np

tic = time.time()

C = nd.zeros((4096, 4096))
for i in range(4096):
    transp = B[[np.arange(4096)], i].reshape(4096,1)
    new_col = nd.dot(A,transp).reshape(4096,)
    C[0:4096, i] = new_col

print(C)
print(time.time() - tic)

```

```

[[-50.474396  -69.83126   -41.304188   ... -19.597466   35.388355
  78.81215    ]
 [-29.70652   -26.495625   96.28163    ...  63.203274   31.081131
  54.01789    ]
 [ 45.610382   28.23838    -1.6904984   ...  45.466324   -6.7074146
 -47.965813   ]
 ...
 [-81.30978   -46.991093  -49.305668   ... -19.005987   -9.508311
 -24.298445   ]
 [-70.84065   113.742485  -85.99479    ... -54.69081    10.467543
 -7.410914    ]
 [ 26.251398  -87.17         63.974815   ...   2.6107101   25.058563
 50.11858     ]]
<NDArray 4096x4096 @cpu(0)>
54.221426486968994

```

In []:

```

# part 4

tic = time.time()

A = nd.normal(0, 1, shape = (4096, 4096))
B = nd.normal(0, 1, shape = (4096, 4096))

C = nd.zeros((4096, 4096))

for j in range(4096):
    b_col = B[[np.arange(4096)], j].reshape(4096,1) #jth col of b

    for i in range(4096):
        a_row_i = A[i].reshape(4096, 1)
        calc = nd.zeros((4096,1))
        test = b_col*a_row_i
        C[i,j] = sum(test)

print(C)
print(time.time() - tic)

```

2. Semidefinite Matrices

Assume that $A \in \mathbb{R}^{m \times n}$ is an arbitrary matrix and that $D \in \mathbb{R}^{n \times n}$ is a diagonal matrix with nonnegative entries.

1. Prove that $B = A D A^{\top}$ is a positive semidefinite matrix.
2. When would it be useful to work with B and when is it better to use A and D ?

part 1

$$B = A D A^{\top}$$

$$A^{\top} B = A^{\top} A D A^{\top} A$$

$$A^{\top} B A = D A^{\top} A A$$

$$D = A^{\top} B A$$

it's given that D is diagonal with nonnegative entries, therefore D is a positive semidefinite matrix.

and so for an arbitrary matrix A , if D can be decomposed as $A^{\top} B A$ and is ≥ 0 , B is also a PSD.

part 2

one might use A and D when trying to check if the eigen values for B are also nonnegative. also, calculating the determinant for D will be easier, especially if B is not diagonal. we may use B for covariance matrix.

3. MXNet on GPUs

1. Install GPU drivers (if needed)
2. Install MXNet on a GPU instance
3. Display `!nvidia-smi`
4. Create a 2×2 matrix on the GPU and print it. See http://d2l.ai/chapter_deep-learning-computation/use-gpu.html for details.

In [1]:

```
import mxnet as mx
```

In [2]:

```
from mxnet import nd
```

In [3]:

```
from mxnet.gluon import nn
```

In [4]:

```
mx.cpu(), mx.gpu(), mx.gpu(1)
```

Out[4]:

```
(cpu(0), gpu(0), gpu(1))
```

In [5]:

```
x = nd.ones((2, 2), ctx = mx.gpu())  
x
```

Out[5]:

```
[[1.  1.]  
 [1.  1.]]  
<NDArray 2x2 @gpu(0)>
```

4. NDArray and NumPy

Your goal is to measure the speed penalty between MXNet Gluon and Python when converting data between both. We are going to do this as follows:

1. Create two Gaussian random matrices A , B of size 4096×4096 in NDArray.
2. Compute a vector $\mathbf{c} \in \mathbb{R}^{4096}$ where $c_i = \|A B_{i \cdot}\|^2$ where \mathbf{c} is a **NumPy** vector.

To see the difference in speed due to Python perform the following two experiments and measure the time:

1. Compute $\|A B_{i \cdot}\|^2$ one at a time and assign its outcome to \mathbf{c}_i directly.
2. Use an intermediate storage vector \mathbf{d} in NDArray for assignments and copy to NumPy at the end.

In [6]:

```
# part 1  
A = nd.normal(0, 1, shape = (4096, 4096))  
B = nd.normal(0, 1, shape = (4096, 4096))
```

```
c = sum((nd.dot(A, B)**2) ** (1/2))
c = c.asnumpy()
c
```

Out[6]:

```
array([4068.9065, 4076.6711, 4124.69 , ..., 4072.759 , 4069.4392,
       4062.4224], dtype=float32)
```

In [9]:

```
# part 2

tic = time.time()

d = nd.zeros((4096, 1))
product = nd.dot(A,B)
for i in range(4096):
    sum_sq = sum(product[0:4096, i]**2)
    d[i] = sum_sq

c_i = d.reshape(4096, 1).asnumpy()

print(c_i)
print(time.time() - tic)
```

```
[16556000.]
[16619247.]
[17013068.]
...
[16587366.]
[16560336.]
[16503276.]
851.4459578990936
```

In [18]:

```
c_i = (d.reshape(4096,1).asnumpy())** (1/2)
c_i
```

Out[18]:

```
array([[4068.9065],
       [4076.6711],
       [4124.69 ],
       ...,
       [4072.759 ],
       [4069.4392],
       [4062.4224]], dtype=float32)
```

5. Memory efficient computation

We want to compute $\$C \leftarrow A \cdot B + C$, where A , B and C are all matrices. Implement this in the most memory efficient manner. Pay attention to the following two things:

1. Do not allocate new memory for the new value of C .
2. Do not allocate new memory for intermediate results if possible.

In [19]:

```
C += nd.dot(A, B)
C
```

Out[19]:

```
[[ -89.2842      73.64483   -25.558414   ...    90.13895    95.803665
  -162.3457   ]
 [ -37.078636   -6.3714905   11.097342   ...    44.199265   -45.428867
   78.536705   ]
 [ 134.2446    214.59921    68.38664    ...   -66.71899   -20.114132
   40.64216    ]
```

```
...
[ -95.55035      41.435562   122.75974   ...   69.42801   -94.71486
  46.208897 ]
[-231.5051     -25.642136   -83.37668   ...   -16.456684   -135.8109
  -83.9735    ]
[  75.27684     -1.9213676    21.20346   ...   81.998985    -7.3830795
  -84.509315  ]]
<NDArray 4096x4096 @cpu(0)>
```

6. Broadcast Operations

In order to perform polynomial fitting we want to compute a design matrix A with

$$A_{ij} = x_i^j$$

Our goal is to implement this **without a single for loop** entirely using vectorization and broadcast. Here $1 \leq j \leq 20$ and $x = \{-10, -9.9, \dots, 10\}$. Implement code that generates such a matrix.

In [20]:

```
x = nd.array(np.arange(-10, 10, 0.1)).reshape(10, 20)
j = nd.array(np.arange(1, 21))
A = x**j
A
```

Out[20]:

```
[[-1.00000000e+01  9.80099945e+01 -9.41192078e+02  8.85292676e+03
 -8.15372812e+04  7.35091875e+05 -6.48477400e+06  5.59581920e+07
 -4.72161280e+08  3.89416294e+09 -3.13810596e+10  2.46990275e+11
 -1.89790670e+12  1.42321134e+13 -1.04106316e+14  7.42510859e+14
 -5.16116234e+15  3.49466736e+16 -2.30389674e+17  1.47808970e+18]
 [-8.00000000e+00  6.24099998e+01 -4.74552032e+02  3.51530371e+03
 -2.53552520e+04  1.77978516e+05 -1.21512812e+06  8.06460250e+06
 -5.19986840e+07  3.25524288e+08 -1.97732659e+09  1.16463319e+10
 -6.64685240e+10  3.67322071e+11 -1.96407892e+12  1.01534516e+13
 -5.07060400e+13  2.44416288e+14 -1.13616609e+15  5.08857954e+15]
 [-6.00000000e+00  3.48100014e+01 -1.95112015e+02  1.05559998e+03
 -5.50731738e+03  2.76806406e+04 -1.33892531e+05  6.22597062e+05
 -2.77990500e+06  1.19042400e+07 -4.88281240e+07  1.91581280e+08
 -7.18019648e+08  2.56666829e+09 -8.73710080e+09  2.82748436e+10
 -8.68351672e+10  2.52599534e+11 -6.94602170e+11  1.80167691e+12]
 [-4.00000000e+00  1.52100010e+01 -5.48719978e+01  1.87416107e+02
 -6.04661682e+02  1.83826562e+03 -5.25233594e+03  1.40640840e+04
 -3.51843750e+04  8.19628047e+04 -1.77147000e+05  3.53814938e+05
 -6.50211000e+05  1.09419012e+06 -1.67725850e+06  2.32830650e+06
 -2.90798000e+06  3.24414975e+06 -3.20649900e+06  2.78218175e+06]
 [-2.00000000e+00  3.60999990e+00 -5.83199978e+00  8.35210133e+00
 -1.04857607e+01  1.13906250e+01 -1.05413494e+01  8.15730476e+00
 -5.15978241e+00  2.59374309e+00 -1.00000000e+00  2.82429427e-01
 -5.49755916e-02  6.78222906e-03 -4.70185274e-04  1.52587891e-05
 -1.71798746e-07  3.87420762e-10 -5.24288153e-14  1.00000029e-20]
 [-3.55271368e-14  1.00000007e-02  8.00000038e-03  8.10000114e-03
  1.02400007e-02  1.56250000e-02  2.79936083e-02  5.76480031e-02
  1.34217739e-01  3.48678321e-01  1.00000000e+00  3.13842916e+00
  1.06993265e+01  3.93737450e+01  1.55568054e+02  6.56840820e+02
  2.95147998e+03  1.40630918e+04  7.08234922e+04  3.75899625e+05]
 [ 2.00000000e+00  4.40999937e+00  1.06480007e+01  2.79840984e+01
  7.96262512e+01  2.44140625e+02  8.03180786e+02  2.82429565e+03
  1.05784551e+04  4.20707383e+04  1.77147000e+05  7.87662500e+05
  3.68934950e+06  1.81633140e+07  9.37959200e+07  5.07094240e+08
  2.86511642e+09  1.68900588e+10  1.03726137e+11  6.62662414e+11]
 [ 4.00000000e+00  1.68099995e+01  7.40879898e+01  3.41880157e+02
  1.64916235e+03  8.30376562e+03  4.35817578e+04  2.38112781e+05
  1.35260600e+06  7.97922800e+06  4.88281240e+07  3.09629280e+08
  2.03255949e+09  1.37994691e+10  9.68069366e+10  7.01137224e+11
  5.23837217e+12  4.03410424e+13  3.19986875e+14  2.61240419e+15]
 [ 6.00000000e+00  3.72099991e+01  2.38327972e+02  1.57529626e+03
  1.07374189e+04  7.54188906e+04  5.45516000e+05  4.06067575e+06
  3.10871080e+07  2.44619440e+08  1.97732659e+09  1.64096788e+10
  1.39740496e+11  1.22045058e+12  1.09263695e+13  1.00225956e+14
  9.41523068e+14  9.05384045e+15  8.90835745e+16  8.96482683e+17]
 [ 8.00000000e+00  6.56100082e+01  5.51367981e+02  4.74583252e+03
  4.18211836e+04  3.77149500e+05  3.47927925e+06  3.28211620e+07
```

```
3.16478432e+08 3.11817062e+09 3.13810596e+10 3.22475655e+11
3.38253002e+12 3.62044059e+13 3.95291543e+14 4.40126666e+15
4.99587160e+16 5.77950934e+17 6.81232885e+18 8.17906469e+19]]
<NDArray 10x20 @cpu(0)>
```