

Homework 1 - Berkeley STAT 157

Handout 1/22/2017, due 1/29/2017 by 4pm in Git by committing to your repository. Please ensure that you add the TA Git account to your repository.

1. Write all code in the notebook.
2. Write all text in the notebook. You can use MathJax to insert math or generic Markdown to insert figures (it's unlikely you'll need the latter).
3. **Execute** the notebook and **save** the results.
4. To be safe, print the notebook as PDF and add it to the repository, too. Your repository should contain two files: homework1.ipynb and homework1.pdf.

The TA will return the corrected and annotated homework back to you via Git (please give `rythei` access to your repository).

```
In [1]: from mxnet import ndarray as nd
```

1. Speedtest for vectorization

Your goal is to measure the speed of linear algebra operations for different levels of vectorization. You need to use `wait_to_read()` on the output to ensure that the result is computed completely, since `NDArray` uses asynchronous computation. Please see http://beta.mxnet.io/api/ndarray/autogen/mxnet.ndarray.NDArray.wait_to_read.html for details.

1. Construct two matrices A and B with Gaussian random entries of size 4096×4096 .
2. Compute $C = AB$ using matrix-matrix operations and report the time.
3. Compute $C = AB$, treating A as a matrix but computing the result for each column of B one at a time. Report the time.
4. Compute $C = AB$, treating A and B as collections of vectors. Report the time.
5. Bonus question - what changes if you execute this on a GPU?

```
In [2]: # part 1
A = nd.normal(0, 1, shape = (4096, 4096))
B = nd.normal(0, 1, shape = (4096, 4096))

import time

# part 2
tic = time.time()
C = nd.dot(A,B)
print(C)
print(time.time() - tic)

[[ 5.1015415 -26.479822  51.007053 ... -7.666253  101.429016
 -101.521935 ]
 [-2.1350682  42.957554 -27.059336 ...  50.63337  29.264576
  85.48195 ]
 [ 179.61961  104.32006 -42.85021 ...  4.7839203 -12.166644
  33.128284 ]
 ...
 [-19.171553 -24.856712  72.80983 ...  36.38448 -28.801601
  8.885335 ]
 [-181.80132  2.513298 -8.9951515 ... -78.88404 -61.0984
 -65.812164 ]
 [ 64.217995  35.151516 -58.099884 ...  39.043835  62.036263
  1.4770873]]
<NDArray 4096x4096 @cpu(0)>
1.947239637374878
```

```
In [5]: # part 3

import numpy as np

tic = time.time()

C = nd.zeros((4096, 4096))
for i in range(4096):
    transp = B[np.arange(4096)], i).reshape(4096,1)
    new_col = nd.dot(A,transp).reshape(4096,)
    C[0:4096, i] = new_col

print(C)
print(time.time() - tic)

[[ 5.1015415 -26.479822  51.007053 ... -7.666253  101.429016
 -101.521935 ]
 [-2.1350682  42.957554 -27.059336 ...  50.63337  29.264576
  85.48195 ]
 [ 179.61961  104.32006 -42.85021 ...  4.7839203 -12.166644
  33.128284 ]
 ...
 [-19.171553 -24.856712  72.80983 ...  36.38448 -28.801601
  8.885335 ]
 [-181.80132  2.513298 -8.9951515 ... -78.88404 -61.0984
 -65.812164 ]
 [ 64.217995  35.151516 -58.099884 ...  39.043835  62.036263
  1.4770873]]
<NDArray 4096x4096 @cpu(0)>
53.83018636703491
```

comment Q1.4

i couldn't find a solution that was efficient enough to compute in time for the deadline, so i did a 100x100 matrix instead. i know the point of the exercise was just to show how slow it is not using efficient vectorization, but i still solved the problem and i hope this will suffice. thanks!

(- :

```
In [6]: # part 4

tic = time.time()

A_100 = nd.normal(0, 1, shape = (100, 100))
B_100 = nd.normal(0, 1, shape = (100, 100))

C_100 = nd.zeros((100, 100))

for j in range(100):
    b_col = B[np.arange(100)], j].reshape(100,1) #jth col of b

    for i in range(100):
        a_row_i = A[i].reshape(100, 1)
        calc = nd.zeros((100,1))
        test = b_col*a_row_i
        C_100[i,j] = sum(test)

print(C_100)
print(time.time() - tic)

[[-12.168565 -4.50321  8.210232 ...  0.20940703 -12.027266
 -2.3722553 ]
 [-18.157333 -0.20229495  1.0201511 ...  0.53903174 -5.057396
  2.2450316 ]
 [-6.4128947  4.3309937  6.790632 ... -3.8496003  16.13936
 -2.2113798 ]
 ...
 [ 15.365839 -7.0519075 -7.716825 ... -18.962404  0.6059242
  0.72697425]
 [ 3.2971926  25.627968 -21.552677 ...  2.0595  0.9457317
  16.829914 ]
 [ 12.690775 -6.101557 -6.250719 ... -0.8239009  8.9878
 -1.5448539 ]]
<NDArray 100x100 @cpu(0)>
47.93808698654175
```

2. Semidefinite Matrices

Assume that $A \in \mathbb{R}^{m \times n}$ is an arbitrary matrix and that $D \in \mathbb{R}^{n \times n}$ is a diagonal matrix with nonnegative entries.

1. Prove that $B = ADA^T$ is a positive semidefinite matrix.
2. When would it be useful to work with B and when is it better to use A and D ?

part 1

$$B = ADA^T$$

$$A^T B = A^T ADA^T$$

$$A^T BA = DA^T A$$

$$D = A^T BA$$

it's given that D is diagonal with nonnegative entries, therefore D is a positive semidefinite matrix.

and so for an arbitrary matrix A , if D can be decomposed as $A^T BA$ and is ≥ 0 , B is also a PSD.

part 2

one might use A and D when trying to check if the eigen values for B are also nonnegative. also, calculating the determinant for D will be easier, especially if B is not diagonal. we may use B for covariance matrix.

3. MXNet on GPUs

1. Install GPU drivers (if needed)
2. Install MXNet on a GPU instance
3. Display `!nvidia-smi`
4. Create a 2×2 matrix on the GPU and print it. See http://d2l.ai/chapter_deep-learning-computation/use-gpu.html for details.

```
In [1]: import mxnet as mx
```

```
In [2]: from mxnet import nd
```

```
In [3]: from mxnet.gluon import nn
```

```
In [4]: mx.cpu(), mx.gpu(), mx.gpu(1)
```

```
Out[4]: (cpu(0), gpu(0), gpu(1))
```

```
In [5]: x = nd.ones((2, 2), ctx = mx.gpu())
x
```

```
Out[5]: [[1. 1.]
 [1. 1.]]
<NDArray 2x2 @gpu(0)>
```

4. NDArray and NumPy

Your goal is to measure the speed penalty between MXNet Gluon and Python when converting data between both. We are going to do this as follows:

1. Create two Gaussian random matrices A, B of size 4096×4096 in `NDArray`.
2. Compute a vector $\mathbf{c} \in \mathbb{R}^{4096}$ where $c_i = \|AB_i\|^2$ where \mathbf{c} is a `NumPy` vector.

To see the difference in speed due to Python perform the following two experiments and measure the time:

1. Compute $\|AB_i\|^2$ one at a time and assign its outcome to \mathbf{c}_i directly.
2. Use an intermediate storage vector \mathbf{d} in `NDArray` for assignments and copy to `NumPy` at the end.

```
In [6]: # part 1
A = nd.normal(0, 1, shape = (4096, 4096))
B = nd.normal(0, 1, shape = (4096, 4096))

c = sum((nd.dot(A, B))**2) ** (1/2)
c = c.asnumpy()
c

Out[6]: array([4068.9065, 4076.6711, 4124.69 , ..., 4072.759 , 4069.4392,
 4062.4224], dtype=float32)
```

```
In [9]: # part 2
```

```
tic = time.time()

d = nd.zeros((4096, 1))
product = nd.dot(A,B)
for i in range(4096):
    sum_sq = sum(product[0:4096, i]**2)
    d[i] = sum_sq

c_i = d.reshape(4096, 1).asnumpy()

print(c_i)
print(time.time() - tic)

[[16556000.]
 [16619247.]
 [17013068.]
 ...
 [16587366.]
 [16560336.]
 [16503276.]]
851.4459578990936
```

```
In [18]: c_i = (d.reshape(4096,1).asnumpy())** (1/2)
c_i
```

```
Out[18]: array([[4068.9065],
 [4076.6711],
 [4124.69 ],
 ...,
 [4072.759 ],
 [4069.4392],
 [4062.4224]], dtype=float32)
```

5. Memory efficient computation

We want to compute $C \leftarrow A \cdot B + C$, where A, B and C are all matrices. Implement this in the most memory efficient manner. Pay attention to the following two things:

1. Do not allocate new memory for the new value of C .
2. Do not allocate new memory for intermediate results if possible.

```
In [19]: C += nd.dot(A, B)
C

Out[19]: [[ -89.2842   73.64483 -25.558414 ...  90.13895   95.803665
 -162.3457 ]
 [-37.078636 -6.3714905  11.097342 ...  44.199265 -45.428867
  78.536705 ]
 [ 134.2446   214.59921   68.38664 ... -66.71899 -20.114132
  40.64216 ]
 ...
 [-95.55035   41.435562  122.75974 ...  69.42801 -94.71486
  46.208897 ]
 [-231.5051 -25.642136 -83.37668 ... -16.456684 -135.8109
  83.9735 ]
 [ -75.27684 -1.9213676  21.20346 ...  81.998985 -7.3830795
 -84.509315 ]]
<NDArray 4096x4096 @cpu(0)>
```

6. Broadcast Operations

In order to perform polynomial fitting we want to compute a design matrix A with

$$A_{ij} = x_i^j$$

Our goal is to implement this **without a single for loop** entirely using vectorization and broadcast. Here $1 \leq j \leq 20$ and $x = \{-10, -9.9, \dots, 10\}$. Implement code that generates such a matrix.

```
In [20]: x = nd.array(np.arange(-10, 10, 0.1)).reshape(10, 20)
j = nd.array(np.arange(1, 21))
A = x**j
A

Out[20]: [[-1.00000000e+01  9.80099945e+01 -9.41192078e+02  8.85292676e+03
 -8.15372812e+04  7.35091875e+05 -6.48477400e+06  5.59581920e+07
 -4.72161280e+08  3.89416289e+09 -3.13810596e+10  2.46990275e+11
 -1.89790670e+12  1.42321134e+13 -1.04106316e+14  7.42510859e+14
 -5.1616234e+15  3.49466736e+16 -2.30389674e+17  1.47808970e+18]
 [-8.00000000e+00  6.20499998e+01 -4.74552032e+02  3.51530371e+03
 -2.53552500e+04  1.77978516e+05 -1.21512812e+06  8.06460250e+06
 -5.19986840e+07  3.25524288e+08 -1.97732659e+09  1.16463319e+10
 -6.64685240e+10  3.67322071e+11 -1.96407892e+12  1.01531451e+13
 -5.07060400e+13  2.44416288e+14 -1.13616609e+15  5.08857954e+15]
 [-6.00000000e+00  3.48100014e+01 -1.95112015e+02  1.05559998e+03
 -5.50731738e+03  2.76806406e+04 -1.33892531e+05  6.22597062e+05
 -2.77990500e+06  1.19042400e+07 -4.88281240e+07  1.91581280e+08
 -7.18019648e+08  2.56666829e+09 -8.73710080e+09  2.82748436e+10
 -8.68351672e+10  2.52599534e+11 -6.94602170e+11  1.80167691e+12]
 [-4.00000000e+00  1.52100010e+01 -5.48719978e+01  1.67460210e+02
 -6.04661682e+02  1.83826562e+03 -5.25233594e+03  1.40640840e+04
 -3.51843750e+04  8.19628047e+04 -1.77147000e+05  3.53814938e+05
 -5.50211000e+05  1.09419012e+06 -1.67725850e+06  2.32830650e+06
 3.68934950e+06  1.24414975e+06 -3.20649900e+06  7.8218175e+06]
 [-2.00000000e+00  3.60999990e+00 -5.8319978e+00  8.35210133e+00
 -1.04857607e+01  1.13906250e+01 -1.05413494e+01  1.81573047e+00
 -5.15978241e+00  2.59374309e+00 -1.00000000e+00  2.82429427e-01
 -5.49755916e-02  6.78222906e-03 -4.70185274e-04  1.52587891e-05
 -1.71798746e-07  3.87420762e-10 -5.24288153e-14  1.00000029e-20]
 [-3.55271368e-14  1.00000007e-02  8.00000038e-03  8.10000114e-03
 1.02400007e-02  1.56250000e-02  2.79936083e-02  5.76480031e-02
 1.34217739e-01  3.48678321e-01  1.00000000e+00  3.13842916e+00
 1.06993265e+01  3.93737450e+01  1.55568054e+02  6.58640820e+02
 2.95147998e+03  1.40630918e+04  7.08234922e+04  3.75899625e+05]
 [ 2.00000000e+00  4.40999993e+01  1.06480007e+01  1.06480007e+01
 7.96262312e+01  2.44140625e+02  8.03180786e+02  2.82429565e+03
 -1.05784551e+04  4.20970793e+04 -1.77147000e+05  3.53814938e+05
 3.68934950e+06  1.24414975e+06 -3.20649900e+06  7.8218175e+06]
 [-2.86511642e+09  1.68900588e+10  1.03726137e+11  6.62662414e+11
 4.00000000e+12  1.68099995e+13  7.40879898e+13  3.41880157e+02
 1.64916235e+03  8.30376562e+03  4.35817578e+04  2.38112781e+05
 1.35260600e+06  7.97922800e+06  4.88281240e+07  3.09629280e+08
 2.03255949e+09  1.37994691e+10  9.68069366e+10  7.01137224e+11
 5.23837217e+12  4.03410424e+13  3.19986875e+14  2.61240419e+15]
 [ 6.00000000e+00  3.72099991e+01  2.38327972e+02  1.57529626e+03
 1.07374189e+04  7.54188906e+04  5.45516000e+05  4.60607575e+06
 3.10871080e+07  2.44619440e+08  1.97732659e+09  1.64096788e+10
 1.39740496e+11  1.22045058e+12  1.09263695e+13  8.00225956e+14
 9.41523068e+14  9.05384045e+15  8.90835745e+16  8.96482683e+17]
```